



MARKETING ANALYTICS IN R: STATISTICAL MODELING

Welcome to the Course! Customer Lifetime Value in CRM

Verena Pflieger

Data Scientist at INWT Statistics



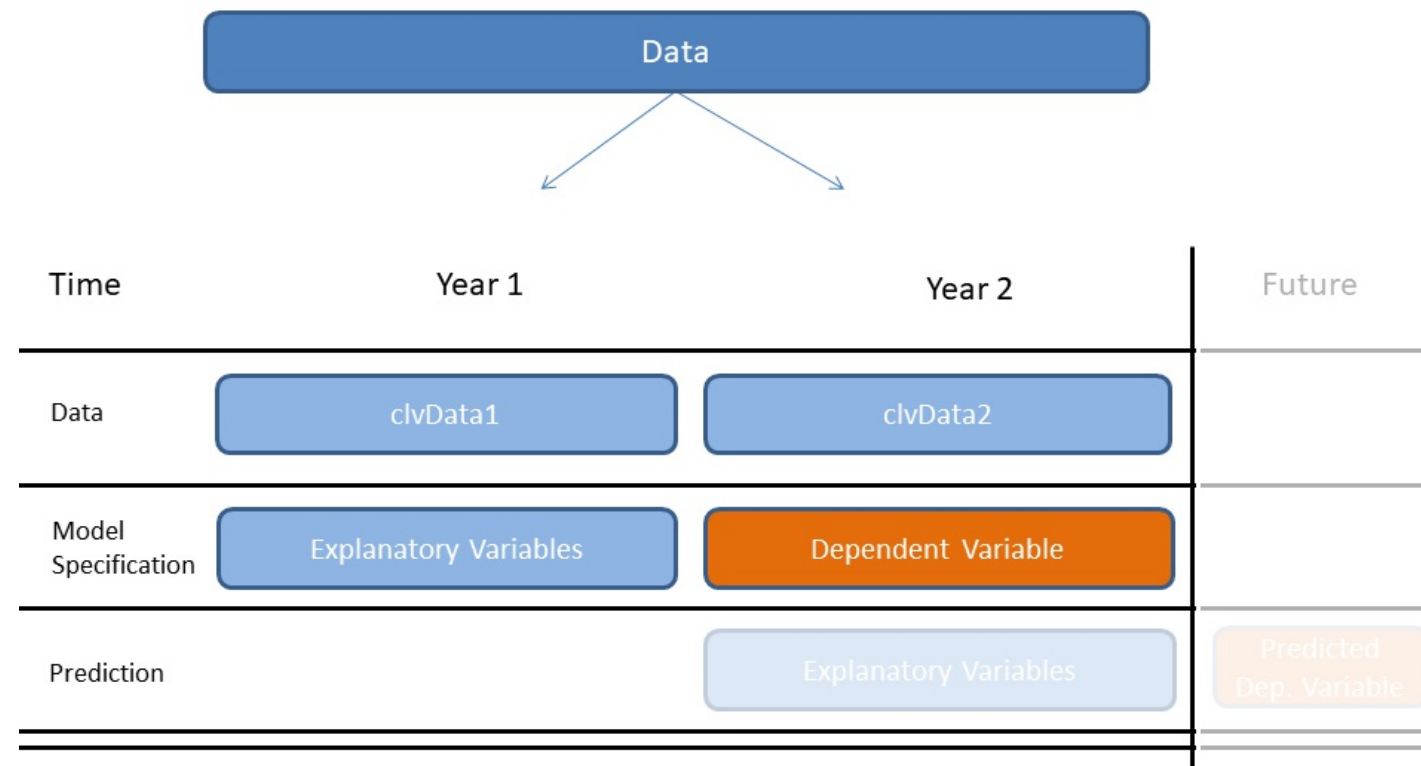
INWT Statistics



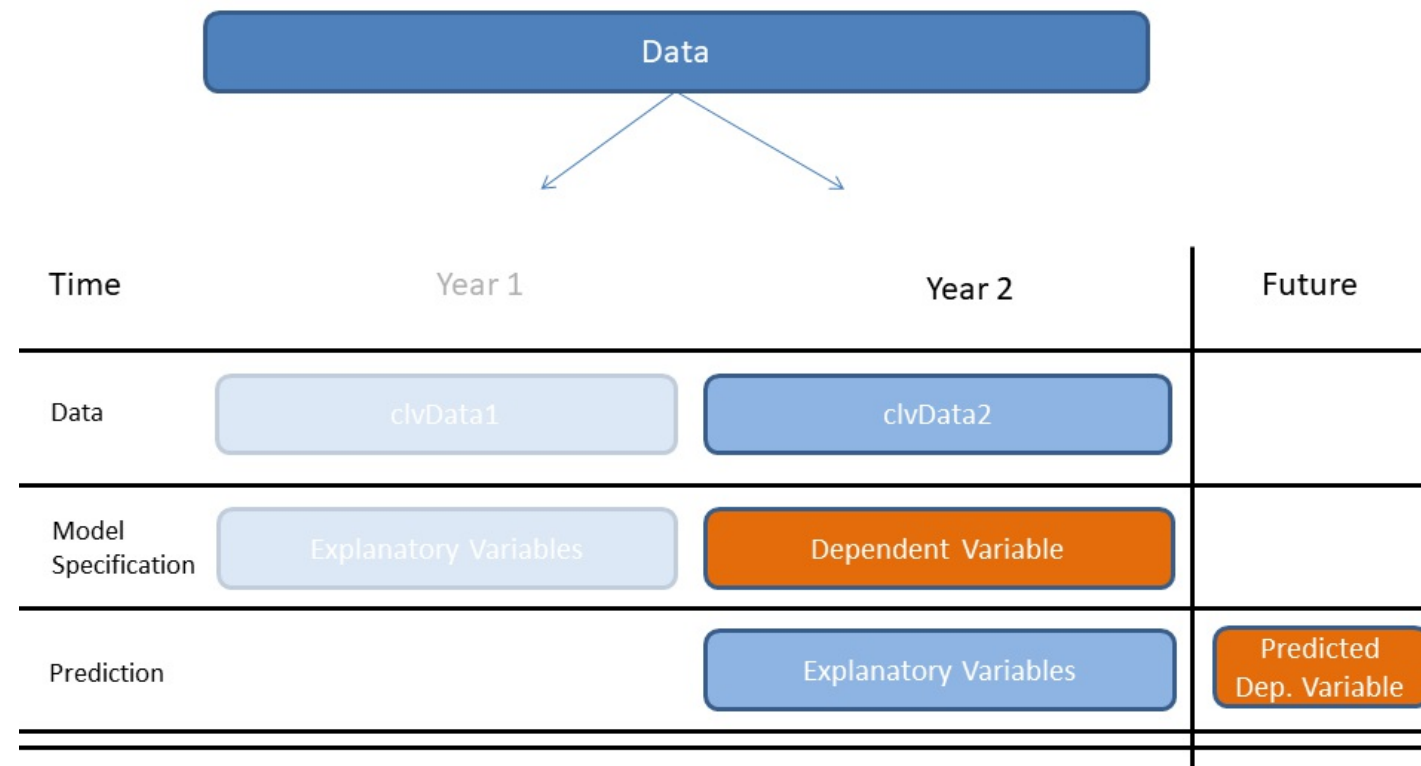
Customer Lifetime Value (CLV)

- predicted future net-profit
- identify promising customers
- prioritize customers according to future margins
- no further customer segmentation

Predicting the Margin of Year 2



Predicting the Future Margin






CLV Data

```
str(clvData1, give.attr = FALSE)
```

```
Classes 'tbl_df', 'tbl' and 'data.frame':    4191 obs. of  15 variables:
 $ customerID      : int  2 3 4 5 6 7 8 9 10 11 ...
 $ nOrders         : int  4 3 12 16 1 2 3 15 16 1 ...
 $ nItems          : int  7 4 25 29 2 8 4 20 18 2 ...
 $ daysSinceLastOrder: int  4 272 12 32 47 19 63 23 75 193 ...
 $ margin          : num  35.8 25.7 63.3 53.7 35.9 ...
 $ returnRatio     : num  0.25 0.44 0.15 0.03 0 0.18 0 0.01 0.02 1 ...
 $ shareOwnBrand   : num  0.67 0.33 0.86 0.96 1 0 0.33 0.53 0.27 0 ...
 $ shareVoucher    : num  0.17 0 0.38 0.17 0 0.86 0.33 0.12 0.6 0 ...
 $ shareSale       : num  0 0.67 0.29 0.33 1 0.14 0 0.12 0.2 1 ...
 $ gender          : chr  "female" "male" "male" "female" ...
 $ age             : int  56 37 32 43 48 31 27 30 50 50 ...
 $ marginPerOrder  : num  8.94 8.58 5.28 3.36 35.85 ...
 $ marginPerItem   : num  5.11 6.43 2.53 1.85 17.93 ...
 $ itemsPerOrder   : num  1.75 1.33 2.08 1.81 2 4 1.33 1.33 1.12 2 ...
 $ futureMargin    : num  57.6 29.7 56.3 58.8 29.3 ...
```

Correlations

```
library(corrplot)
clvData1 %>% select(nOrders, nItems, ...
                    margin, futureMargin) %>% cor() %>% corrplot()
```





MARKETING ANALYTICS IN R: STATISTICAL MODELING

Let's practice!




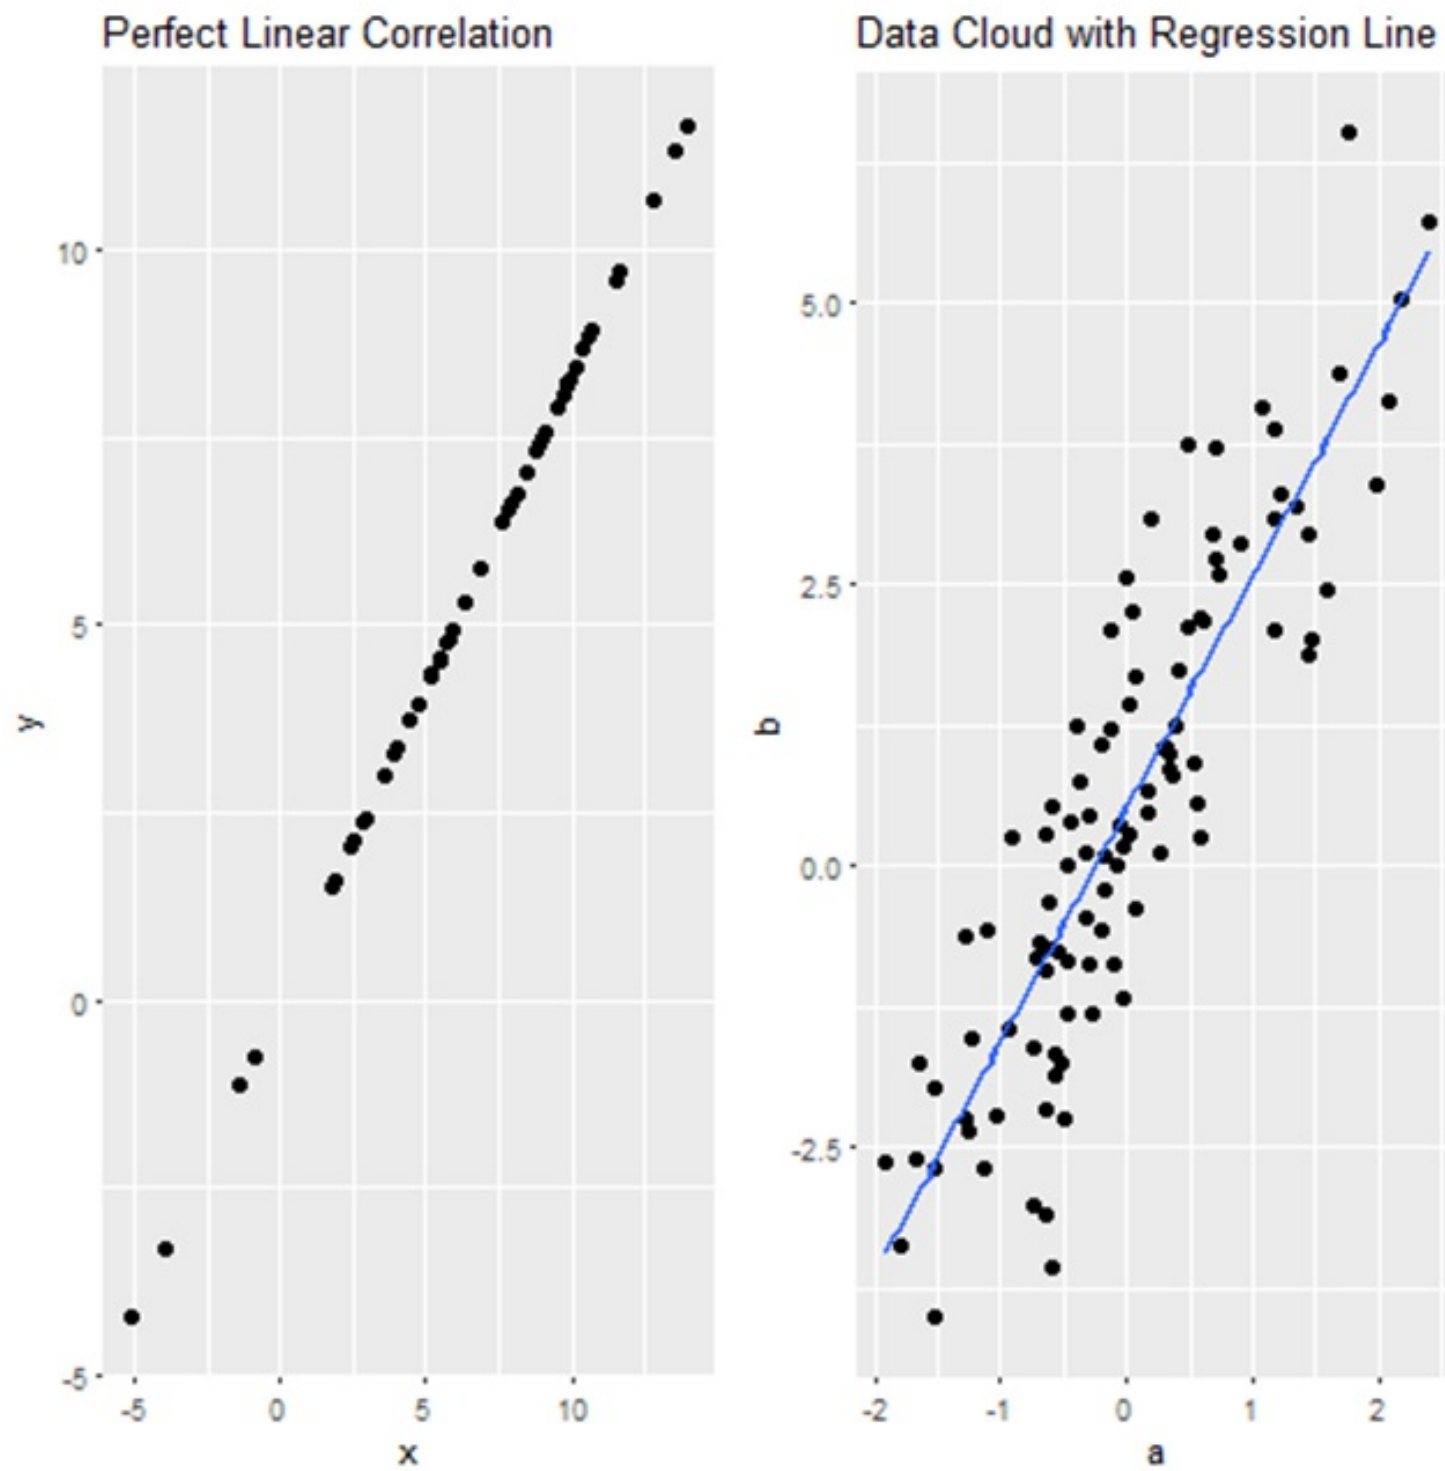
MARKETING ANALYTICS IN R: STATISTICAL MODELING

Simple Linear Regression

Verena Pflieger

Data Scientist at INWT Statistics







Model Specification

```
simpleLM <- lm(futureMargin ~ margin, data = clvData1)
summary(simpleLM)
```

```
Call:
lm(formula = futureMargin ~ margin, data = clvData1)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-56.055	-9.258	0.727	10.060	49.869

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	12.63068	0.49374	25.58	<2e-16 ***
margin	0.64543	0.01467	43.98	<2e-16 ***

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


```
Residual standard error: 14.24 on 4189 degrees of freedom
```

```
Multiple R-squared:  0.3159,    Adjusted R-squared:  0.3158
```

```
F-statistic: 1935 on 1 and 4189 DF,  p-value: < 2.2e-16
```



```
ggplot(clvData1, aes(margin, futureMargin)) +  
  geom_point() +  
  geom_smooth(method = lm, se = FALSE) +  
  xlab("Margin year 1") +  
  ylab("Margin year 2")
```





Assumptions of Simple Linear Regression Model

- Linear relationship between x and y
- No measurement error in x (weak exogeneity)
- Independence of errors
- Expectation of errors is 0
- Constant variance of prediction errors (homoscedasticity)
- Normality of errors



MARKETING ANALYTICS IN R: STATISTICAL MODELING

Time to Practice!



MARKETING ANALYTICS IN R: STATISTICAL MODELING


Multiple Linear Regression

Verena Pflieger

Data Scientist at INWT Statistics




Omitted Variable Bias



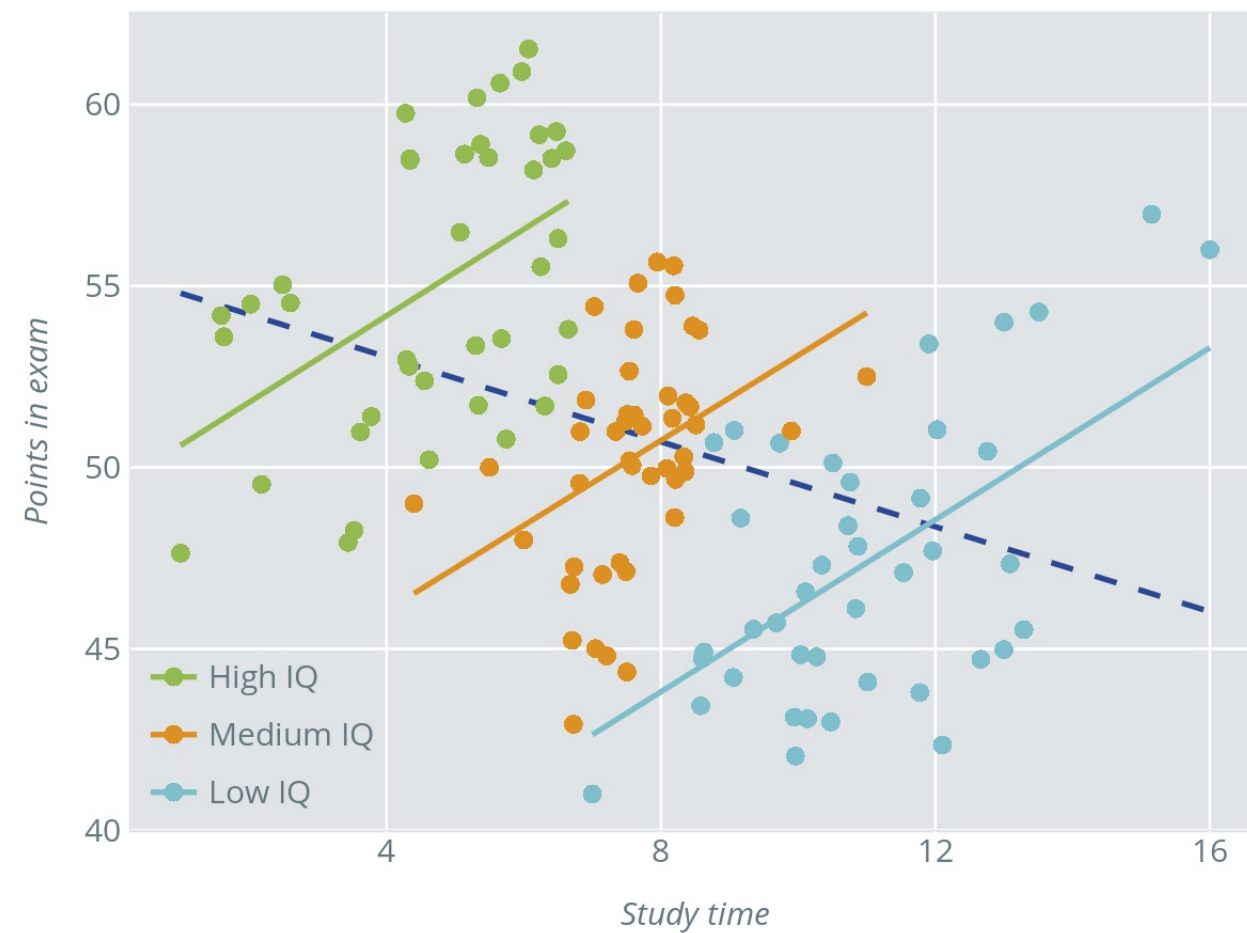


The more Effort, the less Success?





The more Effort, the more Success!





Multiple Linear Regression

```
multipleLM <- lm(futureMargin ~ margin + nOrders + nItems + daysSinceLastOrder +  
                returnRatio + shareOwnBrand + shareVoucher + shareSale +  
                gender + age + marginPerOrder + marginPerItem +  
                itemsPerOrder, data = clvData1)  
  
summary(multipleLM)
```

Call:

```
lm(formula = futureMargin ~ margin + ..., data = clvData1)
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	22.528666	1.435062	15.699	< 2e-16	***
margin	0.402783	0.027298	14.755	< 2e-16	***
nOrders	-0.031825	0.122980	-0.259	0.79581	
...					
itemsPerOrder	0.102576	0.540835	0.190	0.84958	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 13.85 on 4177 degrees of freedom

Multiple R-squared: 0.3547, Adjusted R-squared: 0.3527

F-statistic: 176.6 on 13 and 4177 DF, p-value: < 2.2e-16

Correlation heatmap showing the relationships between 13 variables. The color scale ranges from -1 (dark red) to 1 (dark blue), with 0 being white.

	nOrders	nItems	daysSinceLastOrder	returnRatio	shareOwnBrand	shareVoucher	shareSale	age	marginPerOrder	marginPerItem	itemsPerOrder	margin	futureMargin
nOrders	1.0	0.8	-0.2	-0.2	0.1	0.1	0.1	0.1	-0.4	-0.4	0.1	-0.8	-0.8
nItems	0.8	1.0	-0.2	-0.2	0.1	0.1	0.1	0.1	-0.4	-0.4	0.1	-0.8	-0.8
daysSinceLastOrder	-0.2	-0.2	1.0	-0.6	0.1	0.1	0.1	0.1	0.2	0.2	0.1	-0.2	-0.2
returnRatio	-0.2	-0.2	-0.6	1.0	0.1	0.1	0.1	0.1	0.1	0.1	-0.2	-0.2	-0.2
shareOwnBrand	0.1	0.1	0.1	0.1	1.0	0.6	0.1	0.1	0.1	0.1	0.1	-0.2	-0.2
shareVoucher	0.1	0.1	0.1	0.1	0.6	1.0	0.2	0.1	0.1	0.1	0.1	-0.2	-0.2
shareSale	0.1	0.1	0.1	0.1	0.1	0.2	1.0	0.1	0.1	0.1	0.1	-0.2	-0.2
age	0.1	0.1	0.1	0.1	0.1	0.1	0.1	1.0	0.1	0.1	0.1	-0.2	-0.2
marginPerOrder	-0.4	-0.4	0.2	0.1	0.1	0.1	0.1	0.1	1.0	0.6	0.1	-0.2	-0.2
marginPerItem	-0.4	-0.4	0.2	0.1	0.1	0.1	0.1	0.1	0.6	1.0	-0.2	-0.2	-0.2
itemsPerOrder	0.1	0.1	0.1	-0.2	0.1	0.1	0.1	0.1	0.1	-0.2	1.0	-0.8	-0.8
margin	-0.8	-0.8	-0.2	-0.2	-0.2	-0.2	-0.2	-0.2	-0.2	-0.2	-0.8	1.0	0.6
futureMargin	-0.8	-0.8	-0.2	-0.2	-0.2	-0.2	-0.2	-0.2	-0.2	-0.2	-0.8	0.6	1.0

Variance Inflation Factors

```
library(rms)
vif(multipleLM)
```

margin	nOrders	nItems
3.658257	11.565731	13.141486
daysSinceLastOrder	returnRatio	shareOwnBrand
1.368208	1.311476	1.363515
shareVoucher	shareSale	gendermale
1.181329	1.148697	1.003452
age	marginPerOrder	marginPerItem
1.026513	8.977661	7.782651
itemsPerOrder		
6.657435		

New Model

```
multipleLM2 <- lm(futureMargin ~ margin + nOrders +  
                  daysSinceLastOrder + returnRatio + shareOwnBrand +  
                  shareVoucher + shareSale + gender + age +  
                  marginPerItem + itemsPerOrder,  
                  data = clvData1)
```

```
vif(multipleLM2)
```

margin	nOrders	daysSinceLastOrder
3.561828	2.868060	1.354986
returnRatio	shareOwnBrand	shareVoucher
1.305490	1.353513	1.176411
shareSale	gendermale	age
1.146499	1.003132	1.021518
marginPerItem	itemsPerOrder	
1.686746	1.550524	

Interpretation of Coefficients

```
summary(multipleLM2)
```

Call:

```
lm(formula = futureMargin ~ margin + nOrders + ..., data = clvData1)
```

Residuals:

Min	1Q	Median	3Q	Max
-55.659	-8.827	0.483	9.561	50.118

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	22.798064	1.287806	17.703	< 2e-16	***
margin	0.404200	0.026983	14.980	< 2e-16	***
nOrders	0.220255	0.061347	3.590	0.000334	***
daysSinceLastOrder	-0.017180	0.002675	-6.422	1.49e-10	***
returnRatio	-1.992829	0.601214	-3.315	0.000925	***
shareOwnBrand	7.568686	0.677572	11.170	< 2e-16	***
shareVoucher	-1.750877	0.669017	-2.617	0.008900	**
shareSale	-2.942525	0.691108	-4.258	2.11e-05	***
gendermale	0.203813	0.430136	0.474	0.635643	
age	-0.015158	0.017245	-0.879	0.379462	
marginPerItem	-0.197277	0.051160	-3.856	0.000117	***
itemsPerOrder	-0.270260	0.261458	-1.034	0.301354	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1



MARKETING ANALYTICS IN R: STATISTICAL MODELING

Let's practice!



MARKETING ANALYTICS IN R: STATISTICAL MODELING

Model Validation, Model Fit, and Prediction


Verena Pflieger

Data Scientist at INWT Statistics




Coefficient of Determination R^2


R = 1




R = 1




R = 0.36




R = 0.73



R = 0.34



R = 0.05






R^2 and F-test

```
summary(multipleLM2)
```

```
Residual standard error: 13.87 on 4179 degrees of freedom  
Multiple R-squared:  0.3522,    Adjusted R-squared:  0.3504  
F-statistic: 206.5 on 11 and 4179 DF,  p-value: < 2.2e-16
```

Overfitting





Methods to Avoid Overfitting

- AIC() from stats package
- stepAIC() from MASS package
- out-of-sample model validation
- cross-validation
- ...

```
AIC(multipleLM2)
```

```
[1] 33950.45
```

New Dataset clvData2

```
head(clvData2)
```

```
# A tibble: 6 x 14
```

	customerID	nOrders	nItems	daysSinceLastOrder	margin	returnRatio
	<int>	<int>	<int>	<int>	<dbl>	<dbl>
1	2	16	40	2	57.62	0.18
2	3	1	5	124	29.69	1.00
3	4	15	30	68	56.26	0.16
4	5	23	41	103	58.84	0.03
5	6	2	4	104	29.31	0.00
6	7	6	10	41	35.72	0.06

```
# ... with 8 more variables: shareOwnBrand <dbl>, shareVoucher <dbl>,  
# shareSale <dbl>, gender <chr>, age <int>, marginPerOrder <dbl>,  
# marginPerItem <dbl>, itemsPerOrder <dbl>
```



Prediction

```
predMargin <- predict(multipleLM2,  
                      newdata = clvData2)  
head(predMargin)  
  
      1      2      3      4      5      6  
51.10204 31.63335 51.90008 52.62200 36.65194 33.84383  
  
mean(predMargin, na.rm = TRUE)  
[1] 33.95147
```


Learnings Linear Regression

	Learnings Linear Regression
You have learned...	to predict the future customer lifetime value
	to use a linear regression to model a continuous variable
	that the variables for modelling and prediction have to carry the same names



Learnings from the Model

	Learnings from the Model
You have learned...	that the margin in one year is a good predictor for the margin in the following year
	the longer the time since last order, the smaller the expected margin
	characteristics like gender and age don't seem to play a role for the prediction of margin
	etc...



MARKETING ANALYTICS IN R: STATISTICAL MODELING

Alright, Hands On!